


Clean versus deranged few-shot ICL — 128 trials per task



Modality	Model	Best clean acc. % (n)	Best deranged acc. % (n)	Gap	z-gap	p_{task} at n_{max}
Language	Qwen3-14B	30.0% (64)	16.1% (64)	13.9%	3.48	7.9×10^{-21}
Genome	Evo2-40B	33.2% (64)	16.2% (8)	17.0%	4.84	1.6×10^{-22}
Integer seq.	NextTerm-440M	48.1% (64)	16.7% (64)	31.4%	6.95	1.8×10^{-23}
Image	ImageGPT-large	43.5% (31)	19.2% (1)	24.3%	5.99	1.6×10^{-19}
Time-series	TimesFM-2.5	32.6% (48)	15.7% (4)	16.9%	5.12	9.1×10^{-23}
Protein	ProGen2-base	26.4% (48)	16.5% (48)	9.9%	2.83	3.1×10^{-22}

Best clean and deranged shots are selected independently; z-gap uses the paper's unpaired task-cluster SE. Final column: exact one-sided within-task condition-swap p at the maximum shot count. Each p-value tests a pre-specified modality-specific hypothesis and is unadjusted.